

## Data

Number of classes: 11

Training data set has 7213 images

Validation data set has 1543 images

## Training parameters

Batch size: 128

Image size: (256, 256)

Learning rate: 0.001

Number of epochs: 100

Drop out: 50%

Patience: 10

## Data augmentation

I used *ImageDataGenerator* to augment the images while training my model. This technique helps to prevent overfitting and to generalize the model better. There is total of three transformations added:

- Shear the image by 20%
- Zoom 20% in on the image
- Horizontal flip

## Model summary

### 1. Keras Pretrained Model

I used VGG16 architecture to pretrain my model. I set `include_top = false` to exclude the fully-connected layer and only convolutional and pooling layers are imported. The pretrained model summary is showed as below.

```

Model: "vgg16"
Layer (type)                   Output Shape                  Param #
=====
input_1 (InputLayer)          [(None, 256, 256, 3)]        0
block1_conv1 (Conv2D)         (None, 256, 256, 64)         1792
block1_conv2 (Conv2D)         (None, 256, 256, 64)         36928
block1_pool (MaxPooling2D)    (None, 128, 128, 64)         0
block2_conv1 (Conv2D)         (None, 128, 128, 128)        73856
block2_conv2 (Conv2D)         (None, 128, 128, 128)        147584
block2_pool (MaxPooling2D)    (None, 64, 64, 128)          0
block3_conv1 (Conv2D)         (None, 64, 64, 256)          295168
block3_conv2 (Conv2D)         (None, 64, 64, 256)          590080
block3_conv3 (Conv2D)         (None, 64, 64, 256)          590080
block3_pool (MaxPooling2D)    (None, 32, 32, 256)          0
block4_conv1 (Conv2D)         (None, 32, 32, 512)          1180160
block4_conv2 (Conv2D)         (None, 32, 32, 512)          2359808
block4_conv3 (Conv2D)         (None, 32, 32, 512)          2359808
block4_pool (MaxPooling2D)    (None, 16, 16, 512)          0
block5_conv1 (Conv2D)         (None, 16, 16, 512)          2359808
block5_conv2 (Conv2D)         (None, 16, 16, 512)          2359808
block5_conv3 (Conv2D)         (None, 16, 16, 512)          2359808
block5_pool (MaxPooling2D)    (None, 8, 8, 512)            0
=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0

```

## 2. Sequential model

I added my own layers after pretraining the model. The below model was the one with the best accuracy among different attempts:

Convolution layers:

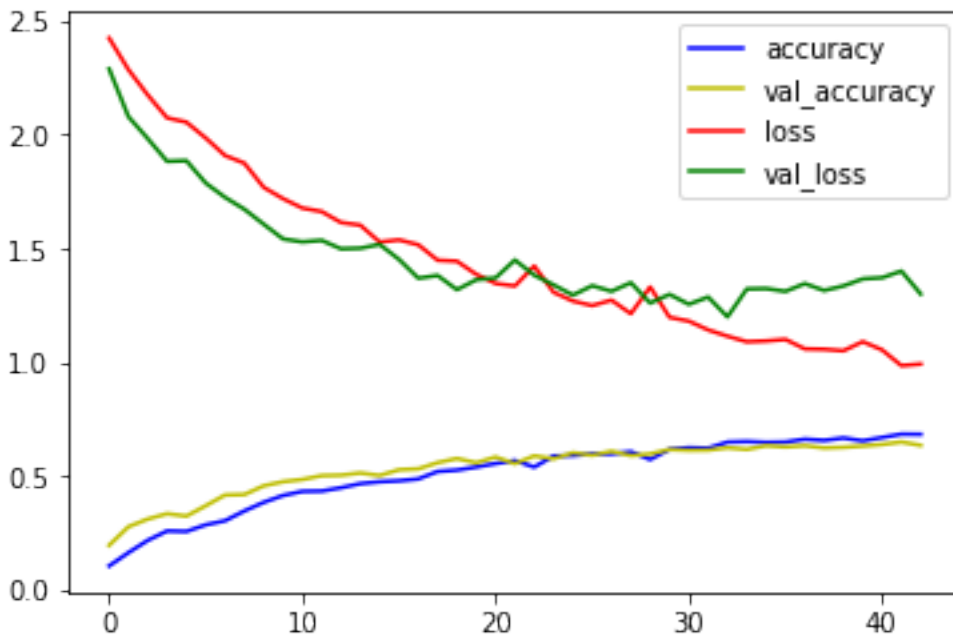
- 1 x convolution layer of 512 channel of 3x3 kernel
- 1 x convolution layer of 1024 channel of 3x3 kernel
- 1 x maxpool layer of 2x2 pool size

RELU activation is added to each layer. After creating the convolution layers, I flatten the vector and pass the data to dense layer. I also used RELU activation for 4 dense layers following with a drop out of 50% for each, and Softmax activation for the last dense layer with a unit of 11 as we have 11 classes to predict:

- 1 x dense layer of 4096 units
- 1 x drop out layer
- 1 x dense layer of 1024 units
- 1 x drop out layer
- 1 x dense layer of 256 units
- 1 x drop out layer
- 1 x dense layer of 32 units
- 1 x drop out layer
- 1 x dense softmax layer of 11 units

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
model (Functional)	(None, 8, 8, 512)	14714688
conv2d (Conv2D)	(None, 6, 6, 512)	2359808
conv2d_1 (Conv2D)	(None, 4, 4, 1024)	4719616
max_pooling2d (MaxPooling2D)	(None, 2, 2, 1024)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 4096)	16781312
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 1024)	4195328
dropout_1 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 256)	262400
dropout_2 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 32)	8224
dropout_3 (Dropout)	(None, 32)	0
dense_4 (Dense)	(None, 11)	363
=====		
Total params: 43,041,739		
Trainable params: 28,327,051		
Non-trainable params: 14,714,688		

## Results



Validation loss: 1.19787

Validation accuracy: 64.87%